

Group 4 - Spotify Song Recommender



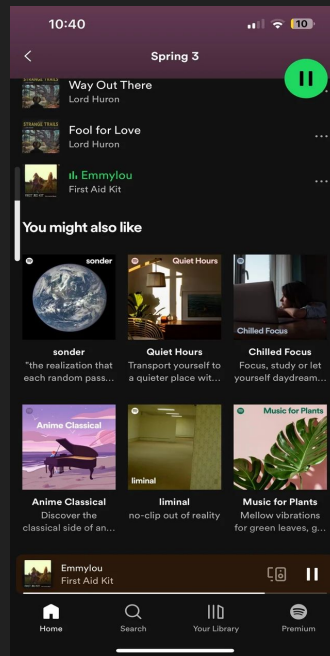
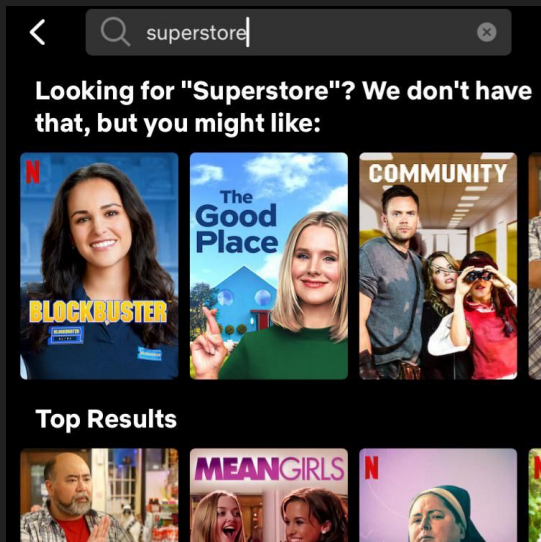
Cagan Abney, Joshua Hale, Jason Cisneros and Raheem Yusuff

[Github](#)

Inspiration

When you are watching Netflix, how do they decide what we should watch next?

How does Spotify decide what song I might want to listen to next? As we all love listening to music and are always looking for that new song we can't stop playing or get out of our head, how can we make finding this song easier?



Dataset

Spotify Tracks Dataset : <https://www.kaggle.com/datasets/maharshipandya/-spotify-tracks-dataset/data>

Columns:

- **Track ID**
- **Track Name**
- **Track Artists**
- **Album Name**
- **Track Genre**
- **Popularity** : between 0 and 100, with 100 being the most popular. Calculated by algorithm on the total number of plays the track has had and how recent those plays are.
- **Audio Features** [song duration, explicitness, danceability, energy, key, loudness, mode (major or minor), speechiness, acousticness, instrumentalness, liveness, valence, tempo, time signature]

Data Engineering

Dropped one row of null values

```
In [5]: 1 df = df.dropna(subset=['album_name'])
        2 df.info()

<class 'pandas.core.frame.DataFrame'>
Index: 113999 entries, 0 to 113999
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   Unnamed: 0             113999 non-null  int64
1   track_id               113999 non-null  object
2   artists                113999 non-null  object
3   album_name             113999 non-null  object
4   track_name             113999 non-null  object
5   popularity              113999 non-null  int64
6   duration_ms            113999 non-null  int64
7   explicit                113999 non-null  bool
8   danceability            113999 non-null  float64
9   energy                 113999 non-null  float64
10  key                    113999 non-null  int64
11  loudness               113999 non-null  float64
12  mode                   113999 non-null  int64
13  speechiness            113999 non-null  float64
14  acousticness           113999 non-null  float64
15  instrumentalness        113999 non-null  float64
16  liveness                113999 non-null  float64
17  valence                 113999 non-null  float64
18  tempo                  113999 non-null  float64
19  time_signature          113999 non-null  int64
20  track_genre             113999 non-null  object
dtypes: bool(1), float64(9), int64(6), object(5)
memory usage: 18.4+ MB
```

Removed all duplicate songs on track ID

```
In [6]: 1 df.drop_duplicates(subset=['track_id'], keep='first', inplace = True)
        2 df.info()

<class 'pandas.core.frame.DataFrame'>
Index: 89740 entries, 0 to 113999
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   Unnamed: 0             89740 non-null  int64
1   track_id               89740 non-null  object
2   artists                89740 non-null  object
3   album_name             89740 non-null  object
4   track_name             89740 non-null  object
5   popularity              89740 non-null  int64
6   duration_ms            89740 non-null  int64
7   explicit                89740 non-null  bool
8   danceability            89740 non-null  float64
9   energy                 89740 non-null  float64
10  key                    89740 non-null  int64
11  loudness               89740 non-null  float64
12  mode                   89740 non-null  int64
13  speechiness            89740 non-null  float64
14  acousticness           89740 non-null  float64
15  instrumentalness        89740 non-null  float64
16  liveness                89740 non-null  float64
17  valence                 89740 non-null  float64
18  tempo                  89740 non-null  float64
19  time_signature          89740 non-null  int64
20  track_genre             89740 non-null  object
dtypes: bool(1), float64(9), int64(6), object(5)
memory usage: 14.5+ MB
```

Dataset Pt. 2

Most Streamed Spotify Songs 2023 Dataset:

<https://www.kaggle.com/datasets/nelgiriyeWithana/top-spotify-songs-2023>

Columns:

- **Track name,**
- **Artist(s) name,**
- **Released year**
- **In spotify playlists**
- **In spotify charts**
- **streams**
- **in_apple_charts**
- **danceability_%**
- **energy_%**
- **acousticness_%**
- **instrumentalness_%**

Data Engineering

This dataset was already very clean, we simply dropped 2 columns that had null values and saved to a new CSV.

```
1 # Check for null values and data types
2 df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 953 entries, 0 to 952
Data columns (total 24 columns):
#   Column                Non-Null Count  Dtype
---  -
0   track_name            953 non-null   object
1   artist(s)_name        953 non-null   object
2   artist_count          953 non-null   int64
3   released_year         953 non-null   int64
4   released_month        953 non-null   int64
5   released_day          953 non-null   int64
6   in_spotify_playlists   953 non-null   int64
7   in_spotify_charts      953 non-null   int64
8   streams               953 non-null   object
9   in_apple_playlists     953 non-null   int64
10  in_apple_charts        953 non-null   int64
11  in_deezer_playlists     953 non-null   object
12  in_deezer_charts       953 non-null   int64
13  in_shazam_charts       903 non-null   object
14  bpm                   953 non-null   int64
15  key                   858 non-null   object
16  mode                  953 non-null   object
17  danceability_%        953 non-null   int64
18  valence_%             953 non-null   int64
19  energy_%              953 non-null   int64
20  acousticness_%        953 non-null   int64
21  instrumentalness_%     953 non-null   int64
22  liveness_%            953 non-null   int64
23  speechiness_%         953 non-null   int64
dtypes: int64(17), object(7)
memory usage: 178.8+ KB
```

```
: 1 # drop columns with null values that we will not use
2 df = df.drop('key', axis=1)
```

```
: 1 df = df.drop('in_shazam_charts', axis=1)
```

Machine Learning: k-Nearest Neighbors (kNN)

- Supervised Machine Learning Model used for Classification or Regression
- kNN was first developed by Evelyn Fix and Joseph Hodges in **1951** in the context of research performed for the US military.
- **How does it work:**
 - Uses proximity to make classifications or predictions about the grouping around an individual data point
 - In order to determine which data points are closest to a given query point, the distance between the query point and the other data points will need to be calculated
 - Based on the k value is how many nearest neighbors the model will predict

kNN Distance Metrics

Euclidean - most commonly used distance measure, measures a straight line between the query point and the other point being measured.

$$d(x,y) = \sqrt{\sum_{i=1}^n (y_i - x_i)^2}$$

Manhattan - referred to as taxicab distance or city block distance as it is commonly visualized with a grid to navigate from one address to another via city streets

$$d(x,y) = \left(\sum_{i=1}^m |x_i - y_i| \right)$$

Minkowski - generalized form of Euclidean and Manhattan distance metrics


$$\left(\sum_{i=1}^n |x_i - y_i| \right)^{1/p}$$

Hamming - typically used with Boolean or string vectors, identifying the points where the vectors do not match.

$$D_H = \left(\sum_{i=1}^k |x_i - y_i| \right)$$


$x=y$	$D=0$
$x \neq y$	$D \neq 1$

kNN Made Simple



kNN: Cosine Similarity

- While the other metrics use distance, Cosine Similarity will use the cosine of the angle between two non-zero vectors



Feature Engineering

Isolated numeric features

```
In [8]: 1 numeric_df['explicit'] = numeric_df['explicit'].astype(float)

In [9]: 1 # Assuming 'df' is your DataFrame
2 int_columns = numeric_df.select_dtypes(include=['int64']).columns
3 numeric_df[int_columns] = numeric_df[int_columns].astype('float64')
4
5 # Verify the data types after conversion
6 print(numeric_df.dtypes)
```

popularity	float64
duration_ms	float64
explicit	float64
danceability	float64
energy	float64
key	float64
loudness	float64
mode	float64
speechiness	float64
acousticness	float64
instrumentalness	float64
liveness	float64
valence	float64
tempo	float64
time_signature	float64
dtype:	object

Encoded Track Genre

```
In [18]: 1 genre_df = pd.get_dummies(df,
2                                     columns = ['track_genre'],
3                                     prefix = ['track_genre'])
4 genre_df.head()
```

Out[18]:

rack_genre_alt-rock	track_genre_alternative	track_genre_ambient	track_genre_anime
False	False	False	False
False	False	False	False
False	False	False	False
False	False	False	False
False	False	False	False

Application Demo: <https://jabney12.pythonanywhere.com/>




Tableau Engineering

Our first tableau dashboard was created with the Most Streamed Spotify Songs 2023 Dataset.

Visualization Titles

- Artists in Most Spotify Playlists
- Spotify vs Apple Charts
- Danceability vs Energy



For the “Artists in Most Spotify Playlists”, the Artist’s bar changes in size based on # of streams per song. We also embedded Spotify to showcase the #1 song.

We added a global filter to change the released years for the Playlist and Charts bar graphs.

By filtering the Danceability vs Energy graph by songs released after the year 2000, there is a clear correlation between the two features.


[Playlist Dashboard](#)

Tableau Engineering

Our second dashboard was created with the Spotify Tracks Dataset

Visualization Titles

- Most Popular Genres
- Genre Instrumentalness vs Genre Acousticness



The Most Popular Genres are shown by bubble size and color, with green being most popular and black being least popular.

Genre Instrumentalness vs Genre Acousticness is shown in a scatter plot with a global Genre filter.

Genre Dashboard

Conclusions

- Taylor Swift is the most popular artists right now
- K-pop is most popular genre as of 2023 on Spotify
- Recommender systems in our everyday life likely use a supervised classification machine learning.
- kNN is a simple classification model good for beginners to learn the complexed world of machine learning.
- Endless amounts of opportunities to optimize machine learning models to achieve your expected outcome.
- Data Science is pretty cool.

Limitations/Bias

Some limitations included our newly introduced machine learning knowledge. With more experience we would be able to use more complex models or better optimize our feature engineering.

Another limitation was the dataset, of course the dataset could not obtain every song in Spotify due to size limitations, so the user can not just input any song.

Limitations also included the data columns, one data set did not have a genre column and one didn't have a time column (Released day, month, year).

Bias also ties into the song limitation, since the number and choice of songs were limited, there is a bias against smaller, less known artists, as well as older songs. Especially since one dataset is based on 2023 popularity

Future Work

- Use other classification machine learning models if KNN truly is the best for a recommendation system.
 - try other distance metrics to see how different our recommendations would be
- Adding more user customization to the application
 - Allow for minor errors in user input (mismatch capitalization or no space)
 - input multiple songs user likes
 - input only an Artist the user likes
 - select what distance metric you would like model to use
 - Auto finish or text box suggestion
- Add more songs

Works Cited

- Algolia. "Cosine Similarity: What Is It and How Does It Enable Effective and Profitable Recommendations?" Algolia Blog, www.algolia.com/blog/ai/cosine-similarity-what-is-it-and-how-does-it-enable-effective-and-profitable-recommendations/.
- IBM. "K-Nearest Neighbors (KNN)." IBM, www.ibm.com/topics/knn.
- Kaggle. "Spotify Tracks Dataset." Kaggle, www.kaggle.com/datasets/maharshipandya/-spotify-tracks-dataset/data.
- Kaggle. "Top Spotify Songs 2023." Kaggle, www.kaggle.com/datasets/nelgiriwewithana/top-spotify-songs-2023.
- "K-Nearest Neighbors (KNN)." Elastic, www.elastic.co/what-is/knn.
- Scikit-learn. "sklearn.neighbors.NearestNeighbors." Scikit-learn Documentation, scikit-learn.org/stable/modules/generated/sklearn.neighbors.NearestNeighbors.html.